

Multiple Partition Regression Analysis

Rhys Hawkins and Malcolm Sambridge

November 9, 2021

This tutorial describes how to analyse discontinuous data using Monte-Carlo Markov-Chain based regression.

Pre-requisites

This tutorial is for the Python version of the `rjmc` library. The examples rely on the `Matplotlib` library for plotting. The versions used in the development of this tutorial are as follows:

- Python 2.7.1
- `rjmc` 0.1.0
- `Matplotlib` 1.1.0

The Data

Similar to the single partition tutorial we will use a non-trivial synthetic dataset with added noise. The difference this time is that we will add a series of discontinuities or step functions and a sign change.

The base function that is used is an exponentially increasing sine wave over the domain $0 \dots 10$, ie:

$$y = \text{stepsign}(x) \times e^{\frac{x}{3}} \sin \frac{2x}{3} + \text{step}(x) \quad (1)$$

Where the `step` and `stepsign` functions are defined as follows:

$$\text{step}(x) = \begin{cases} 15 & x < 2.5 \\ -20 & 2.5 \leq x < 5 \\ 0 & \text{otherwise} \end{cases} \quad (2)$$

$$\text{stepsign}(x) = \begin{cases} -1 & x < 2.5 \\ 1 & 2.5 \leq x < 5 \\ -1 & \text{otherwise} \end{cases} \quad (3)$$

The actual dataset unevenly (though with fairly good coverage) samples this function and adds some Gaussian noise and these values are save to an ASCII text file. A plot of the synthetic data points with the true function is shown in Figure 1.




Figure 1: The Synthetic Data

Loading the Data

For those unfamiliar with Python, we discuss briefly the loading of data from a simple ASCII text file. The code for loading the data and doing a simple plot is shown in Listing 1.

Listing 1: Loading the Data

```

1  #!/usr/bin/env python3
2  #-- coding: utf-8 --
3  #
4  # Import the libraries we will need for analysis and plotting.
5  #
6  import rjmc
7  import matplotlib
8  import matplotlib.pyplot
9
10 #
11 # Open our data file which consists of one (x, y) coordinate per line
12 # separated by whitespace
13 #
14 f = open('data.txt', 'r')
15 lines = f.readlines()
16
17 x = []
18 y = []
19
20 for line in lines:
21     columns = line.split()
22     x.append(float(columns[0]))
23     y.append(float(columns[1]))
24
25 f.close()
26
27 fig = matplotlib.pyplot.figure()
28
29 matplotlib.pyplot.plot(x, y, 'ko')
30
31 fig.savefig('ch1-loading.pdf', format='PDF')
32 matplotlib.pyplot.show()
33

```

The ASCII file contains an x, y pair per line separated by a space. On line 12, we use the built in function open to open the file. On lines 15 and 16 we initialize 2 list that will contain the x and y coordinates in the file. On line 18 we loop through the lines in the file and add each x, y pair to the 2 separate lists.

From line 26 onwards, we plot the data using the matplotlib library and save the plot to a PDF file. The plot resulting from this script can be seen in Figure 2.




Figure 2: The Plot of the Data

Running the default analysis

For performing a regression analysis on a continuous dataset the function is called `regression_part1d`. The parameters for this function are mostly the same as the `regression_single1d` function with the only additional required parameter being `pd`.

The complete list of parameters for this function are as follows with default values shown where applicable:

dataset The dataset object to run the analysis on. This is an `rjmc.dataset1d` object which wraps the x and y vectors you load from the file and includes individual point noise values. This is the only parameter which doesn't have a default value.

pd The standard deviation for the perturbation of partition boundaries.

burnin = 10000 The number of initial samples to throw away.

total = 50000 The total number of samples to use for the analysis.

max_order = 1 The maximum order of polynomial to use to fit the data.

xsamples = 100 The number of points to sample along the x direction for the curve.

ysamples = 100 The number of points to sample along the y directory for the statistics such as mode, median and confidence intervals. This is the number of bins for the histograms in the y direction.

confidence_interval = 0.95 The confidence interval to use for minimum and maximum confidence intervals. This should be a value between 0 and 1.

For this analysis we are only going to use the default values and the listing is shown in Listing 2.

Listing 2: Running the Default Analysis

```

1  #!/usr/bin/env python3
2  # ← coding: utf-8 → ←
3  #
4  # Import the libraries we will need for analysis and plotting.
5  #
6  import rjmc
7  import matplotlib
8  import matplotlib.pyplot
9
10 #
11 # Open our data file which consists of one (x, y) coordinate per line
12 # separated by whitespace
13 #
14 f = open('data.txt', 'r')
15 lines = f.readlines()
16
17 x = []
18 y = []
19
20 for line in lines:
21     columns = line.split()
22     x.append(float(columns[0]))
23     y.append(float(columns[1]))
24
25 f.close()
26
27 #
28 # Set our x range
29 #
30 xmin = 0.0
31 xmax = 10.0
32
33 #
34 # Estimate our error standard deviation
35 #
36 sigma = 3.0
37 n = [sigma] * len(x)
38
39 #
40 # Create the rjmc dataset
41 #
42 data = rjmc.datasetid(x, y, n)
43
44 #
45 # Specify the standard deviation for the move partition
46 #
47 pd = 1.0
48
49 #
50 # Run an analysis with reduced max order to only allow linear
51

```

```

52 # segments
53 #
54 burnin = 10000
55 total = 50000
56 max_partitions = 10
57 max_order = 1
58
59 results = rjmc regression_part1d(data,
60                                pd,
61                                burnin,
62                                total,
63                                max_partitions,
64                                max_order)
65
66 #
67 # Retrieve the mean curve for plotting
68 #
69 xc = results.x()
70 meancurve = results.mean()
71
72 #
73 # Retrieve the partition location and count information
74 #
75 partlocation = results.partition_location_histogram()
76 partcount = results.partitions()
77
78 #
79 # Plot the data with black crosses and the mean with a red line
80 #
81 fig = matplotlib.pyplot.figure(1)
82
83 a = matplotlib.pyplot.subplot(211)
84
85 a.plot(x, y, 'ko', xc, meancurve, 'r-')
86 a.set_xlim(xmin, xmax)
87
88 b = matplotlib.pyplot.subplot(212)
89 b.bar(xc, partlocation, xc[1] - xc[0])
90 b.set_xlim(xmin, xmax)
91
92 fig.savefig('ch2-analyse.pdf', format='PDF')
93
94 fig = matplotlib.pyplot.figure(2)
95
96 a = matplotlib.pyplot.subplot(111)
97 a.hist(partcount, bins=5, range=(0, 5), align='left')
98
99 fig.savefig('ch2-analyse-partcount.pdf', format='PDF')
100
101 matplotlib.pyplot.show()

```

The preamble (lines 1 ... 24) consists of loading the file as in the previous section.

An important part of the analysis is estimating the error in the data. This is specified as a error value per data point and can be thought of a weighting as to how well the fit will attempt to fit an individual point. If the value is low, then the fit will be tight and conversely if the value is high then the fit will be loose. On lines 35 and 36 we set a value of 3.0 for all data points. Use this value for now, but try other values greater than 0.0 to see the effect.

On line 41 we construct the `dataset1d` object from the `x`, `y` and `n` lists we created. These lists must be the same length.

On line 46 we set the value for the `pd` parameter. With partition modelling, a variable number of discontinuities are trialed at random locations along the `x`-axis. As part of the fitting process, the locations of these discontinuities are perturbed by an amount determined by sampling from a normal distribution with a standard deviation given by the `pd` parameter. We will discuss how to choose this value latter, but for now a useful rule of thumb would be to set it to 5 to 10 percent of the range of the data. So in this case our range is 10, so a `pd` of 1 would seem a good first approximation.

On line 57 we run the analysis with this `dataset1d` object. The

`regression_single1d` function returns a `resultset1d` object which contains various results and diagnostics about the analysis. For this simple analysis we simply take the x sampling coordinates and the mean of the fits. And plot the mean with the original data points to see how representative the mean is. This plot is shown in Figure 3.

In the plot, rather than just plotting the mean of the fits, we have also plotted the histogram of the location of the partition boundaries which indicates the most likely location of the discontinuities. As can be seen in the figure, the histogram highlights the artificial discontinuities created at $x = 2.5$ and 5 .

We have also plotted the histogram of the number of partitions as shown in Figure 4. From this we can infer the likelihood of the number of the partitions in the data.

It should be noted here that there are two discontinuities in function and one less prominent change in gradient in the real data. Looking at the plot of the number of partitions, we see strong support for 4 partitions, as the algorithm has identified the four segments in the original curve.




Figure 3: The Default Analysis Plot




Figure 4: The Partition Count Histogram

Increasing the maximum order of the polynomial

Partition modelling and allowing higher order polynomials are in conflict to some degree. A higher order polynomial can fit a discontinuity to some degree and therefore obviate the need to create an extra partition.

To show this, we can increase the maximum order allowed for the polynomials in each partition to 5. This means that the fitting procedure may use up to a 5th order polynomial instead of line segments. The script to do this is shown in Listing 3.

Listing 3: Increasing the maximum order of the polynomials

```

1  #!/usr/bin/env python3
2  # -*- coding: utf-8 -*-
3  #
4  # Import the libraries we will need for analysis and plotting.
5  #
6  import rjmc
7  import matplotlib
8  import matplotlib.pyplot
9
10 #
11 # Open our data file which consists of one (x, y) coordinate per line
12 # separated by whitespace
13 #
14 f = open('data.txt', 'r')
15 lines = f.readlines()
16
17 x = []
18 y = []
19
20 for line in lines:
21     columns = line.split()
22     x.append(float(columns[0]))
23     y.append(float(columns[1]))
24
25 f.close()
26
27 #
28 # Set our x range
29 #
30

```

```

31 xmin = 0.0
32 xmax = 10.0
33
34 #
35 # Estimate our error standard deviation
36 #
37 sigma = 3.0
38 n = [sigma] * len(x)
39
40 #
41 # Create the rjcmc dataset
42 #
43 data = rjcmc.dataset1d(x, y, n)
44
45 #
46 # Specify the standard deviation for the move partition
47 #
48 pd = 1.0
49
50 #
51 # Run the default analysis
52 #
53 results = rjcmc.regression_part1d(data, pd)
54
55 #
56 # Retrieve the mean curve for plotting
57 #
58 xc = results.x()
59 meancurve = results.mean()
60
61 #
62 # Retrieve the partition location and count information
63 #
64 partlocation = results.partition_location_histogram()
65 partcount = results.partitions()
66
67 #
68 # Plot the data with black crosses and the mean with a red line
69 #
70 fig = matplotlib.pyplot.figure(1)
71
72 a = matplotlib.pyplot.subplot(211)
73
74 a.plot(x, y, 'ko', xc, meancurve, 'r-')
75 a.set_xlim(xmin, xmax)
76
77 b = matplotlib.pyplot.subplot(212)
78 b.bar(xc, partlocation, xc[1] - xc[0])
79 b.set_xlim(xmin, xmax)
80
81 fig.savefig('ch3-order.pdf', format='PDF')
82
83 fig = matplotlib.pyplot.figure(2)
84
85 a = matplotlib.pyplot.subplot(111)
86 a.hist(partcount, bins=5, range=(0, 5), align='left')
87
88 fig.savefig('ch3-orderpartcount.pdf', format='PDF')
89
90 matplotlib.pyplot.show()

```

The results of running this script can be seen in Figures 5 and 6. As can be seen, the fit is more smooth than the previous analysis and supports a fewer number of partitions. In particular we can see that there is some support for no discontinuities, but it is more likely that there are 1 or 2 which matches our artificially created boundaries (the discontinuity at $x = 5$ is not that large so it can be accommodated by a higher order polynomial). Try changing the maximum allowed order to, say, 2, i.e. allow up to a quadratic, and examine what happens to the Partition Count Histogram.

Confidence

So far we have only plotted the mean of the fits, however this gives only a glimpse as to how the data is being fitted. One way to look at the fitting is to take samples of the fitting curves. The script to do this is showing in Listing 4.




Figure 5: The Order Analysis Plot




Figure 6: The Partition Count Histogram

Listing 4: Sampling the Fitting

```

1  #!/usr/bin/env python3
2  #-- coding: utf-8 --
3  #
4  # Import the libraries we will need for analysis and plotting.
5  #
6  import rjmc
7  import matplotlib
8  import matplotlib.pyplot
9  from mpl_toolkits.mplot3d import Axes3D
10
11 #
12 # Open our data file which consists of one (x, y) coordinator per line
13 # separated by whitespace
14 #
15 f = open('data.txt', 'r')
16 lines = f.readlines()
17
18 x = []
19 y = []
20
21 for line in lines:
22     columns = line.split()
23     x.append(float(columns[0]))
24     y.append(float(columns[1]))
25
26 f.close()
27
28 #
29 # Estimate our error standard deviation
30 #
31 #
32 sigma = 3.0
33 n = [sigma] * len(x)
34
35 #
36 # Create the rjmc dataset
37 #
38 data = rjmc.dataset1d(x, y, n)
39
40 #
41 # Set our x range
42 #
43 xmin = 0.0
44 xmax = 10.0
45
46 #
47 # This is our callback function which samples the curves generated
48 # during the analysis
49 #
50 sample_x = None
51 sample_curves = []
52 sample_i = 0
53 sample_rate = 250
54 def sampler_cb(x, y):
55     global sample_x, sample_curves, sample_i, sample_rate
56
57     if sample_i == 0:
58         sample_x = x
59
60     if sample_i % sample_rate == 0:
61         sample_curves.append(y)
62
63     sample_i = sample_i + 1
64
65 #
66 # Run a series of analyses with varying maximum allowed order
67 #
68 #
69 # Specify the standard deviation for the move partition
70 #
71 pd = 1.0
72
73 #
74 # Run an analysis with reduced max order to only allow linear
75 # segments
76 #
77 burnin = 10000
78 total = 50000
79 max_partitions = 10
80 max_order = 3
81
82 results = rjmc.regression_part1d_sampled(data,
83                                         sampler_cb,
84                                         pd,
85                                         burnin,
86                                         total,
87                                         max_partitions,
88                                         max_order)
89
90 #
91 # Plot the data with black crosses, the sample curves as faint lines, and
92 # the mean as a red line
93 #
94 fig = matplotlib.pyplot.figure(1)
95 ax = fig.add_subplot(111)
96
97 yc = 0.5
98 yalpha = 1.0 / ((1.0 - yc) * float(len(sample_curves)))
99 for sy in sample_curves:
100

```

```

101     ax.plot(sample_x, sy,
102             color = str(y_c),
103             alpha = y_alpha,
104             linestyle = '-',
105             linewidth = 10)
106
107     ax.plot(results.x(), results.mean(), 'r-')
108     ax.plot(x, y, 'ko')
109     ax.set_xlim(xmin, xmax)
110
111     fig.savefig('ch4-confidence.pdf', format='PDF')
112
113     fig = matplotlib.pyplot.figure(2)
114     ax = fig.add_subplot(111)
115
116     ax.plot(results.x(), results.mean(), 'r-')
117     ax.plot(x, y, 'ko')
118     ax.plot(results.x(), results.credible_min(), 'b:')
119     ax.plot(results.x(), results.credible_max(), 'b:')
120
121     ax.set_xlim(xmin, xmax)
122
123     fig.savefig('ch4-confidenceintervals.pdf', format='PDF')
124
125     matplotlib.pyplot.show()

```

The resulting plot is shown in Figure 7. This plot shows the mean as well as 200 curves sampled from the fitting process overplotted with transparency so that the darker regions are where the true curve is more likely to be. What should be noted is that there are large spikes about the discontinuities and this is to be expected as the fitting process will trial fitting curves across these discontinuities which will produce high order curves which will overshoot.

Another feature to notice is that at approximately $x = 4.5$, the curves tend to take 2 alternate paths. This type of feature is likely due to whether a partition boundary is located at approximately $x = 5$ or not.

An alternative look at the breadth of curves is to plot the confidence intervals that are generated during the fitting process. This result is shown in Figure 8.

The confidence interval curves are shown in blue dashed lines. What these curves represent is 95 percent of the curves generated during the fitting process were contained within these lines. Similarly to the sampled curves, there is some overshoot near severe discontinuities. Lowering the maximum order can sometimes alleviate these.




Figure 7: Sampled Curves from the Fitting




Figure 8: Confidence Intervals from the Curve Fitting